

Dissection: Issue Map v1

Sourceless Political Claim Dissection v4 run against No Party Politics · Issue Map v1

September 16, 2026

Pre-retrieval run. Fetching disabled; no claim requiring external retrieval reaches SUPPORTED or CONTRADICTED. Two platform texts held from earlier in session were used.

Contents

- Run header 3**
- Headline 4**
- Stage 0 — Artifact intake 5**
 - Denatured claims. 5
- Stage 2 — Atomization 6**
 - Platform claims adjudicated against held text. 6
- Stage 4 — Domain audit 8**
 - F1 — Evidentiary asymmetry by party · MATERIAL 8
 - F2 — Draft framed as adopted · MATERIAL 8
 - F3 — Speech cited as platform · MATERIAL 8
 - F4 — Citation-stacking to an interested T4 origin · MATERIAL 9
 - F5 — Probable source misattribution · MATERIAL 9
 - F6 — No margin of error anywhere · MATERIAL 9
 - F7 — Definition and denominator risk on a headline number · MATERIAL 10
 - F8 — Missing baseline on a rate claim · MATERIAL 10
 - F9 — Unstated selection rule · MATERIAL 10
 - F10 — Header overclaim · MINOR 10
 - F11 — Compression distortion · MINOR 11
 - F12 — Stale figures carried without in-body ageing · MINOR 11
 - F13 — Lookalike-domain check required · HYPOTHESIS, NOT A FINDING 11
 - F14 — Date precision · MINOR 11
 - F15 — About the object, not the map · INFORMATIONAL 11
- Stage 5 — Argument audit 12**
- Stage 6 — Plain-language fidelity gate 13**
- Stage 7 — Influence-operation patterns 14**
- Stage 8 — Red-team pass 15**
- Stage 9 — Retrieval plan 17**
- Verdicts 19**
- Fix list, ordered 20**
- Machine-readable output 21**

Run header

Procedure: claim-dissection@4.0.1-runtime-ascii. Object under analysis: *What Each Party Says It Stands For, and What Its Voters Actually Think* — No Party Politics · Issue Map v1, dated September 16, 2026.

Mode: pre-retrieval with two held artifacts. No --fetch flag was passed, so links in the Issue Map were not dereferenced during this run. Under §2.3 that means **no claim requiring external retrieval can reach SUPPORTED or CONTRADICTED**. That is the designed behavior, not a failure.

Two exceptions, declared: the 2024 Republican and Democratic platform texts were retrieved from presidency.ucsb.edu earlier in this same session (2026-09-16 03:27 UTC) — the identical URLs the Issue Map cites as [GOP24] and [DEM24] — and are held locally. They were used. Every platform claim below was checked against that verbatim text. Survey claims were not, and are all citation_pending.

Markers. No <<<ARTIFACT>>> markers were supplied. Per §2.2 the whole pasted blob is treated as one artifact, a1, modality text, role object_under_analysis. Operator instruction was the line preceding it.

Injection: none detected. No hidden text, no counterfeit system block, no verdict-steering language, no nested or forged markers.

Refusals fired: none. No private individuals, no unnamed persons in media, no geolocation request, no authorship attribution, no persuasive-content request.

Conformance: false. No weights hash, no seed, no source snapshot id. Stage 8 red-team ran in-context rather than as a separate invocation (red_team_contaminated = true). Content hashes below are SHA-256 of the local extracted text, not of the served HTTP response; they verify what was read, not what the server sent.

Headline

The Issue Map's characterizations of both party platforms **hold up under verbatim check**. Every GOP and Democratic platform claim I could test against the held text was accurate, several near-verbatim. That is the substantive finding and it is a good result for the document.

The problems are not in what it says about the platforms. They are in **how it sources one party versus the other**, in an **entirely unmarked margin of error**, and in **one probable citation swap** on an AI figure.

Three things to fix before publication, in order of severity:

1. **Six of fifteen Democratic stances are sourced to third-party summaries when the adopted platform text says the same thing.** All fifteen Republican stances cite the platform. This is the single most consequential finding in the run.
2. **Three voter-ID percentages are sourced to a White House press release.** An interested party's press release cannot carry a survey figure.
3. **An AI statistic is attributed to an advocacy pollster but matches a Pew pair exactly.** If it is Pew's, the document's own caveat is discounting a figure that does not need discounting.

Stage 0 — Artifact intake

Field	Value
artifact_id	a1
modality	text
role	object_under_analysis
tier	T4
tier rationale	Self-published analytical document. Under Rule 2 it does not inherit the tier of the sources it cites. Its 46 reference entries are pointers, not evidence.
extraction	direct paste, confidence 1.00, not truncated

Held evidence artifacts, used:

id	source	retrieved	sha256 (extract, first 16)	tier
e1	2024 Republican Party Platform, presidency.ucsb.edu	2026-09-16 03:27 UTC	2140dad5fe7d29d0	T1-equivalent
e2	2024 Democratic Party Platform, presidency.ucsb.edu	2026-09-16 03:27 UTC	55ccfd525108c086	T1-equivalent

Tier note, stated because the schema does not cover it cleanly. A party platform is not law and is not official government data, so it fits no tier as written. But for an ATTRIBUTIVE claim about its own contents — “the platform says X” — a document published by the issuing body is primary, on the same principle that makes official transcripts T1. I have tiered e1 and e2 as T1-equivalent **for attributive claims about platform content only**. For any claim about the world they are T4. This is a schema gap in v4 and should be closed in v5 with an explicit “issuing-body primary document” tier.

Denatured claims

D1. The 2024 Republican and Democratic platforms state the positions this document attributes to them, across fifteen issue areas.

D2. Named nonpartisan pollsters measured the partisan opinion figures this document reports, on the dates given.

D3. Every line in the document points to a source listed at the bottom.

D4. The named issues will or will not carry into the 2028 presidential cycle. (*Predictive. Labeled as judgment by the author.*)

Timeframe: 2021–August 2026. Actors: RNC, DNC, named pollsters. Jurisdiction: U.S. federal. Source of authority: party platforms and survey research.

Stage 2 — Atomization

Approximately 150 discrete assertions are present. Atomizing all of them would produce a document longer than the one under analysis, so this run adjudicates a **prioritized subset**: every claim I could test against held evidence, plus every claim carrying a failure pattern, plus the load-bearing document-level claims. The remaining ~120 survey figures are dispositioned as a class in the JSON (`unverified_items`) with a single verdict and a shared retrieval plan. That is an engineering decision and it is stated rather than hidden.

Platform claims adjudicated against held text

Republican, all SUPPORTED against e1 (confidence 0.95, pinpoint):

Issue Map claim	Platform text
"the platform does not mention the Affordable Care Act by name"	0 occurrences of "Affordable Care"; 0 of "Obamacare"
"The 2024 platform uses [abortion] once"	exactly 1 occurrence
"no change to the retirement age"	core promises: "NO CUTS, INCLUDING NO CHANGES TO THE RETIREMENT AGE"
"open limited federal land for home building"	"open limited portions of Federal Lands to allow for new home construction"
"revoke China's Most Favored Nation status"	verbatim
"block Chinese vehicle imports"	"preventing the importation of Chinese vehicles"
"cancel the electric vehicle mandate"	core promises: "CANCEL THE ELECTRIC VEHICLE MANDATE"
"defund schools teaching critical race theory or what it calls gender ideology"	core promises: "CUT FEDERAL FUNDING FOR ANY SCHOOL PUSHING CRITICAL RACE THEORY, RADICAL GENDER IDEOLOGY"
"close the federal Department of Education"	"close the Department of Education in Washington, D.C. and send it back to the States"
"build an Iron Dome missile defense shield"	core promises: "BUILD A GREAT IRON DOME MISSILE DEFENSE SHIELD OVER OUR ENTIRE COUNTRY"
"listed among the 20 core promises" (guns)	core promises: "THE RIGHT TO KEEP AND BEAR ARMS"
"confront prosecutors it labels Marxist"	"stand up to Marxist Prosecutors"
"reassert federal control over Washington DC"	"reassert greater Federal Control over Washington, DC"
"same day voting"	core promises: "SECURE OUR ELECTIONS, INCLUDING SAME DAY VOTING, VOTER IDENTIFICATION, PAPER..."
"Repeal the Biden AI executive order"	"We will repeal Joe Biden's dangerous Executive Order that hinders AI Innovation"

Issue Map claim	Platform text
"Oppose a central bank digital currency; defend Bitcoin mining and self-custody"	1 occurrence each of "central bank digital", "Bitcoin", "self-custody"

Democratic, all SUPPORTED against e2 (confidence 0.95, pinpoint):

Issue Map claim	Platform text
"help build or renovate 2 million homes"	verbatim
"the platform's own estimate is \$2,500 a year" (tariffs)	"tariffs will slow growth... raising taxes on middle-class families by \$2,500 a year"
"Medicare negotiate at least 50 drug prices a year"	"add at least 50 drugs a year to that list, lowering prices for 500 drugs this decade"
"add dental, vision and hearing to traditional Medicare"	"expand traditional Medicare coverage to include dental, vision, and hearing services"
"Pass the John Lewis Voting Rights Advancement Act and Freedom to Vote Act"	"pass... the John R. Lewis Voting Rights Advancement Act and the Freedom to Vote Act"
"the Equality Act, which would bar discrimination against LGBTQ+ people"	"pass the Equality Act to codify protections for LGBTQI+ Americans"
universal background checks; assault weapons and magazine ban; safe storage; national red flag law; end gun industry liability protection	all five present verbatim in one passage
"supporting long-term undocumented residents including Dreamers"	"Dreamers" ×5
codify Roe; contraception; IVF; medication abortion	"Roe" ×8, "contraception" ×7, "IVF" ×6

Net: 25 of 25 testable platform claims SUPPORTED. No contradictions found.

Stage 4 — Domain audit

F1 — Evidentiary asymmetry by party · MATERIAL

All fifteen Republican stances cite the platform text [GOP24]. Six of fifteen Democratic stances cite something else:

Issue	Democratic stance sourced to	What it actually is
4. Immigration	[NTD_DEMPLAT]	news summary of the platform
7. Abortion	[DNC_DRAFT]	DNC release on the platform draft
8. Guns	[RELOAD_DEMGUNS]	gun-policy outlet's summary
10. Elections	[DNC_DRAFT]	DNC release on the platform draft
14. Foreign policy	[BRI]	third-party platform-comparison page
15. LGBTQ	[ABC_S0TU24]	State of the Union liveblog

I checked all six against the adopted text. **All six planks are in it**, in several cases in wording closer to the Issue Map's own summary than the substitute source is. The guns section is the clearest case: the Issue Map lists five specific gun planks from a secondary outlet, and the adopted platform states all five consecutively in a single passage.

So the substitution is not necessary. Its effect is that one party is characterized from its own words and the other from other people's summaries of its words. In a document whose entire premise is neutral side-by-side mapping, that is structural, not clerical.

Fix: replace all six with pinpoint [DEM24] quotes. The text is already in hand.

F2 — Draft framed as adopted · MATERIAL

[DNC_DRAFT] is titled "DNC Releases 2024 Party Platform **Draft**." It is the source of record for the abortion and elections stances. A draft is not the adopted document — this is the introduced-framed-as-enacted pattern applied to a platform. The claims survive because the adopted text agrees, but the citation does not survive. The adopted platform was formally adopted 2024-08-19, which the document's own source list already records.

F3 — Speech cited as platform · MATERIAL

[ABC_S0TU24] is a State of the Union liveblog. A presidential address is not a party platform. It sits under a heading that reads "Official stance — Democratic Party (2024 platform)," which mislabels the source class. The adopted platform has a dedicated LGBTQI+ section naming the Equality Act.

F4 — Citation-stacking to an interested T4 origin · MATERIAL

Issue 10 sources three voter-ID percentages — Republicans 98%, independents 84%, Democrats 67% — to [WH_VOTERID], a White House press release archived under the title “VOTER ID OVERWHELMINGLY POPULAR WITH LITERALLY EVERYONE EXCEPT...”. The document’s own caveat concedes the numbers are “relayed through a White House release.”

A press release from a party to the dispute is T4 and cannot carry a survey figure, whatever the underlying poll says. Separately, 98% is an extreme value for any survey item and is worth confirming directly. The citable artifact is the Gallup release (October 2024), which the document already cites elsewhere as [GALLUP_VOTE24].

Fix: cite Gallup directly, or drop the three figures.

F5 — Probable source misattribution · MATERIAL

Issue 13 reports “61% not confident in the government’s ability to regulate AI” (Republicans) and “74%” (Democrats), attributed to [NBC_AIPI26].

The identical pair — Republicans 61%, Democrats 74%, same item — appears in Pew’s *Americans and AI 2026* report, fieldwork February 17-23, 2026, pages 17-18. **This project already holds that citation** in `research/ecosystem.md`.

Two different pollsters returning the same two integers on the same question is improbable. Either NBC is relaying Pew and the label is wrong, or the figures have been crossed with the neighboring AIPI items.

This matters beyond bookkeeping. The document flags AIPI as advocacy-aligned and tells the reader to discount it. If the numbers are Pew’s, the reader is being told to discount a figure that does not need discounting — a caveat misapplied is its own small distortion.

Decisive artifact: Pew *Americans and AI 2026* PDF, pp. 17-18, against the NBC piece.

F6 — No margin of error anywhere · MATERIAL

Not one figure in the document carries a margin of error or a sample size. The midterm table then ranks one-point differences as findings:

| AI | 22% | 23% | 51% | **R+1** Economy | 37% | 36% | 25% | **D+1**

At typical survey sample sizes a one-point gap is not distinguishable from zero. Presenting it in a column headed “Edge,” ordered from largest R advantage to largest D advantage, invites the reader to treat it as directional. The large gaps in the same table (Crime R+14, Healthcare D+14, Environment D+14) are almost certainly real; putting them in the same column as R+1 borrows their credibility for the noise.

The table’s arithmetic is internally consistent — I checked all twelve rows — so this is a presentation problem, not a computation error.

Fix: publish n and MOE for every figure, and collapse sub-MOE gaps to “tied.”

F7 — Definition and denominator risk on a headline number · MATERIAL

“Border Patrol encounters at the U.S.-Mexico border fell to 237,538 in fiscal 2025, the lowest since 1970.”

“Encounters” as a published series combines Title 8 apprehensions with expulsions and entered routine use after 2020. The 1970 comparison point is an apprehensions series. A 55-year “lowest since” claim requires one consistent metric across the whole span; mixing them is the wrong-denominator pattern.

The underlying point may well survive — Border Patrol apprehensions are a continuous series and the claim may be true on that basis. But as written the metric is not pinned.

Decisive artifact: CBP’s own nationwide encounters and apprehensions time series (T1).

F8 — Missing baseline on a rate claim · MATERIAL

“KFF found the average enrollee’s premium more than doubled and about 9% of last year’s enrollees are now uninsured.”

When a premium subsidy lapses, the net-of-subsidy premium and the gross premium move very differently. The sentence does not say which one doubled, and “more than doubled” is a dramatic claim that lands differently depending on the answer. The 9% figure needs its denominator stated: 9% of last year’s enrollees, or 9% of the current marketplace population.

F9 — Unstated selection rule · MATERIAL

“These are the three most relevant measured positions per group.”

No rule is given for “most relevant.” Choosing three items out of a large survey instrument is exactly where framing enters a document like this, and no procedural protection is offered. This is a cherry-picking exposure, not an accusation — I found no instance where the selection looked slanted, and several where it cut against the obvious narrative (Issue 4 leads with the Republican figure that has moved *toward* more enforcement, Issue 5 leads with a Republican reversal).

Fix: state the rule and apply it uniformly. “Highest salience, largest partisan gap, most recent fieldwork” would do.

F10 — Header overclaim · MINOR

“Every line points to a source at the bottom.”

The fifteen “2028 note” lines carry no source by design, and Issue 6’s “Both 2024 platforms pledge no benefit cuts. They differ on how to fund the programs” makes a claim about both platforms while citing only [GOP24]. **Fix:** “Every factual line...” and add [DEM24] to that sentence.

F11 — Compression distortion · MINOR

The platform reads “promote homeownership through Tax Incentives **and** support for first-time buyers” — two separate items. The Issue Map renders “tax incentives for first-time buyers,” narrowing a general incentive into a first-time-buyer one.

F12 — Stale figures carried without in-body ageing · MINOR

Guns: newest partisan data is October 2024, background checks 2021. Climate “major threat”: 2023. PRRI abortion: 2024 fieldwork. The caveat lines disclose this, but the figures appear undated in the body where a reader will quote them. **Fix:** inline the year on any figure older than twelve months.

F13 — Lookalike-domain check required · HYPOTHESIS, NOT A FINDING

[ABC_S0TU24] and [ABC_ABORT26] use abcnews.com. ABC News’ canonical domain is abcnews.go.com. This may be a legitimate redirect or a transcription slip. **It cannot be resolved with fetching off** and is recorded as a retrieval target, not a finding. Do not act on it until checked.

F14 — Date precision · MINOR

“Enhanced ACA premium credits expired Jan 1, 2026.” The statutory lapse was the end of plan-year 2025; January 1 is when higher premiums took effect. Both framings appear in this project’s own research memo. Pick one, footnote the other.

F15 — About the object, not the map · INFORMATIONAL

The Democratic platform attaches the same “\$2,500 a year” figure to two different policies — a national sales tax in one passage and tariffs in another. The Issue Map cites it correctly for tariffs. But a reader who opens the platform will find one number doing two jobs, which is worth a footnote in a document whose purpose is to help people audit platforms.

Stage 5 — Argument audit

Reconstruction. Stated premises: party platforms state official positions; nonpartisan surveys measure what partisans believe; the two can be compared. Unstated premise: a platform plank and a survey item are commensurable units — that “what the party says” and “what its voters think” are the same kind of object, differing only in who is speaking. They are not. A platform is a negotiated document written once by a committee; a survey item is a response distribution. Comparing them is useful, which is why the document works, but the comparison is looser than the side-by-side layout implies. Conclusion: readers can see where party and base diverge.

Findings:

- *Smuggled presupposition*, material — the “Edge” column presupposes that a one-point difference has a direction. Span: “AI | 22% | 23% | 51% | R+1.”
- *Base-rate / precision neglect*, material — see F6.
- *Cherry-picking exposure*, material — see F9. Exposure, not instance.
- *Vague quantifier*, minor — “More than 8 in 10” (Republicans) sits opposite “5% to 10%” (Democrats) in Issue 9. Two different precisions for the two sides of one comparison; the range is presented without saying it spans two questions.

Credit where due, and it is substantial. The document does several things the method specifically looks for and rarely finds:

- It attributes contested framing to the party that used it rather than adopting it: “confront prosecutors **it labels** Marxist,” “**what it calls** gender ideology,” “the platform **frames this as** keeping men out.” That is correct neutral practice.
- It separates a source’s interpretation from the underlying text: “NPR read the GOP platform’s ‘same day voting’ language as a call to end early voting; **the platform text itself does not spell that out.**” That is the link-fidelity check, performed unprompted.
- It flags its own advocacy-aligned pollsters (Data for Progress, AI Policy Institute) at point of use, not in a footnote.
- It writes “No verified figure gathered for this group” eleven times instead of filling gaps with plausible numbers. Issue 6 and Issue 12 have no partisan data at all and say so. That is the single most important habit in this kind of work.
- It warns against cross-pollster comparison where the question wording differs (Issue 15).

No straw-manning was found in either direction.

Stage 6 — Plain-language fidelity gate

Not triggered for this output. This dissection does not render T0 legal text for the reader; it names provisions and vote records without paraphrasing statutory language.

It should have triggered for the Issue Map. Issue 10 renders bill text:

“The House passed the SAVE America Act on Feb 11, 2026: photo ID for in-person and mail voting, plus DHS citizenship checks of voter rolls.”

That is a summary of statutory content sourced to a fact-checking site rather than to the bill text, and it is unverified in this run. Before publication it needs the full gate: extract the operative provision verbatim with pinpoint citation, preserve the modal verbs exactly, carry every exception, and run the negation-and-scope check. “Photo ID for in-person and mail voting” is precisely the kind of compression where a dropped “only,” “unless,” or “except” changes the meaning. The bill number and Congress number are also missing — required, since bill numbers repeat across Congresses.

Stage 7 — Influence-operation patterns

None to report. This is an analytical document by a known author, not circulating content. No narrative reuse, no coordinated timing, no T4-origin amplification pattern is assessable from a single static artifact, and none is alleged.

One structural note, offered as design guidance rather than a finding: several issues in this map sit on top of grievances that this project has already documented as targets of foreign amplification — AI data centers and electricity bills, tariffs, immigration. A public-facing issue map is a natural surface for that kind of narrative to attach to. Publishing the margin of error and the selection rule (F6, F9) is the cheapest available hardening, because the most common attack on a document like this is selective quotation of a single number.

Stage 8 — Red-team pass

`red_team_contaminated = true` — this ran in-context, not as a separate invocation, so it is not an independent check.

Where this analysis reached past its evidence, caught and corrected mid-run. Two of my three suspected contradictions dissolved when I actually read the text:

1. I flagged “Pass the **John Lewis** Voting Rights Advancement Act” as absent from the adopted platform. A grep for “John Lewis” returned zero. The platform says “**John R. Lewis**” — with the middle initial. My search was wrong, not the document. Had I shipped on the first result I would have published a false CONTRADICTED against a claim that is verbatim accurate.
2. I flagged “reassert federal control over Washington DC” as an interpretation added on top of the core promise “REBUILD OUR CITIES, INCLUDING WASHINGTON DC.” The crime chapter says “Republicans will **reassert greater Federal Control over Washington, DC.**” Near-verbatim. My hypothesis was built from the wrong passage.
3. I also briefly suspected my own prior deliverable had misquoted the GOP platform’s AI heading, because “Artificial Intelligence” returned zero matches. The archived text contains a soft line break — “Artificial Inte lligence” — an extraction artifact in the APP HTML. The heading is correct. Worth knowing for anyone else grepping that file.

Recording these is the point. Three plausible findings, three retractions on contact with the source. That ratio is normal and it is why Stage 3 requires candidate referents to carry `is_hypothesis = true`.

Strongest counter-evidence to the headline finding (F1). The steelman for sourcing Democratic stances to secondary summaries: the adopted Democratic platform is roughly 43,000 words against the Republican platform’s 5,400, and much of it is retrospective narration of the Biden record rather than forward commitments. Extracting a crisp forward-looking plank from it is genuinely harder, and a summary that has already done that work is a reasonable shortcut under time pressure. The asymmetry may be an artifact of document length, not of intent.

That defense is real but it does not survive the test: I extracted all six planks from the adopted text in a single pass, and the guns planks appear consecutively in one paragraph. The work was not hard. The finding stands at material.

Tier inflation found. Three instances: [WH_VOTERID] used as a survey source (F4); [DNC_DRAFT] used as adopted-platform text (F2); [ABC_S0TU24] used as platform text (F3). Each treats a document *about* a source as the source. That is Rule 2, and it is the most common failure in the map.

Tier-confidence mismatches. The document assigns no confidence levels, so no mismatch is computable. But its own hedging is uneven: AIPI and Data for Progress are flagged at point of use, while the White House press release is flagged only in a caveat line beneath the figures it supports. The stronger flag belongs on the weaker source.

Where the tier does not justify my own confidence. The 25 platform verdicts rest on a tier assignment I had to construct because the schema does not cover party platforms. If a reviewer rejects the T1-equivalent call, those 25 verdicts fall to T4 and cap at UNSUPPORTED. I think the call is right and I have stated the reasoning, but the whole SUPPORTED column depends on it.

Symmetry check — required, one line. Reverse the parties: F1 would fire identically, at the same severity, with the same fix. The check is easier here than usual because the two material sourcing findings already land on opposite sides — F1 concerns Democratic stances sourced to third parties, F4 concerns a Republican-administration press release used as a poll source. Neither was softened. The directional finding I would most expect to have gone soft, F5, cuts against a figure that makes both parties look similar, so no partisan gradient is available there either.

Stage 9 — Retrieval plan

Decisive artifact: *Pew, Americans and AI 2026*, pp. 17–18. One fetch resolves F5, and F5 is the only finding that could mean a published number is attached to the wrong pollster.

Primary sources, priority order:

1. *Pew Americans and AI 2026* PDF, pp. 17–18 — resolves F5.
2. Congress.gov for the SAVE America Act, with Congress number — resolves the Stage 6 gate on Issue 10.
3. House Clerk roll-call record for the February 11, 2026 vote — confirms passage, date, and party split.
4. CBP nationwide encounters and apprehensions time series — resolves F7.
5. Gallup, October 2024 voter-ID release — replaces [WH_VOTERID], resolves F4.
6. KFF March 2026 tracking poll methodology page — resolves F8 baseline and supplies n and MOE.
7. Pew methodology pages for every cited survey — resolves F6 across the document.
8. Supreme Court opinion, February 20, 2026 tariff decision — confirms the Issue 2 context claim.

Queries, literal, targeting the claims as written:

1. Pew "Americans and AI 2026" "not confident" government regulate AI Republicans Democrats
2. "SAVE America Act" roll call February 11 2026 clerk.house.gov
3. CBP nationwide Border Patrol apprehensions fiscal year 2025 237,538
4. Gallup October 2024 voter identification photo ID party Republicans Democrats independents
5. KFF March 2026 ACA enhanced premium tax credits average premium increase methodology
6. NBC News June 2026 AI Policy Institute poll regulation confidence

Branch outcomes:

If found	Then
Pew pp. 17–18 carries 61/74	F5 confirmed; relabel to Pew, remove the AIPI discount on that line
NBC independently reports 61/74 from AIPI	F5 dissolves; note the coincidence and keep both citations
CBP series shows a consistent apprehensions metric to 1970	F7 downgrades to minor; pin the metric name in the sentence
CBP series shows “encounters” begins post-2020	F7 stands; rewrite the comparison or drop “since 1970”
Gallup October 2024 shows 98/84/67	F4 stands as a sourcing fix only; swap the citation, keep the numbers
Gallup shows different figures	F4 escalates to disqualifying; remove the three figures

If found	Then
SAVE Act text matches the one-line summary	Stage 6 gate passes; add bill and Congress number
SAVE Act text carries exceptions not summarized	Stage 6 gate fails; render verbatim

Verdicts

Per claim class:

Class	Count	Verdict	Confidence	Basis
Platform stance claims, tested	25	SUPPORTED	0.95	held T1-equivalent text, pinpoint
Platform stance claims, untested	~8	UNSUPPORTED	0.30	not reached this pass
Survey figures	~120	UNSUPPORTED , citation_pending	0.30	fetch off; all are retrieval targets
"2028 note" judgments	15	UNVERIFIABLE	n/a	predictive; correctly labeled by author
"Every line points to a source"	1	MISLEADING_CONTEXT	0.90	minor; see F10
Legislative claims (SAVE Act, SCOTUS tariff ruling)	2	UNSUPPORTED, citation_pending	0.30	T0/T1 resolvable, not fetched

Post-level rollup: MIXED, confidence 0.55, advisory. Load-bearing claims split: the platform layer is supported against held text, the survey layer is entirely unverified in this run, and nine material findings fired. Not downgraded to MISLEADING_CONTEXT, because that requires the load-bearing set to be otherwise uniformly supported, which it is not.

What MIXED means here, plainly: the document is not wrong about the platforms. It is under-sourced on the numbers, asymmetrically sourced between the parties, and imprecise about measurement error. Those are all fixable without changing a single substantive conclusion.

Fix list, ordered

1. Replace six Democratic stance citations with pinpoint [DEM24] quotes. Text is already extracted and in this project. *(F1, F2, F3)*
2. Verify the 61/74 AI pair against Pew pp. 17–18 and relabel if it is Pew’s. *(F5)*
3. Replace [WH_VOTERID] with the Gallup original, or cut the three voter-ID figures. *(F4)*
4. Add n and margin of error to every figure; collapse sub-MOE gaps in the midterm table to “tied.” *(F6)*
5. Pin the border metric or drop “lowest since 1970.” *(F7)*
6. State whether the doubled ACA premium is net or gross, and give the 9% denominator. *(F8)*
7. State the three-item selection rule and apply it uniformly. *(F9)*
8. Run the SAVE America Act summary through the fidelity gate; add bill and Congress number. *(Stage 6)*
9. “Every factual line points to a source”; add [DEM24] to the Issue 6 both-platforms sentence. *(F10)*
10. Give each 2028 note a mechanism and a base rate instead of a bare likelihood, or relabel the column “Author’s judgment, unscored.” *(D4)*
11. Inline the fieldwork year on any figure older than twelve months. *(F12)*
12. Check the abcnews.com domain before acting on it. *(F13)*
13. Split the GOP housing tax-incentive plank back into its two items. *(F11)*
14. Footnote the platform’s reuse of “\$2,500” for two policies. *(F15)*

Machine-readable output

```
{
  "schema_version": "4",
  "contract_version": "sotp/0.1.0-draft",
  "prompt_version": "claim-dissection@4.0.1-runtime-ascii",
  "run": {
    "mode": "pre-retrieval",
    "fetch_enabled": false,
    "conformant": false,
    "nonconformance_reasons": [
      "no weights_hash available",
      "no seed available",
      "no source_snapshot_id available",
      "Stage 8 red-team ran in-context, not as a separate invocation",
      "content hashes are of local extracted text, not of the served HTTP response",
      "atomization is a prioritized subset; ~120 survey figures dispositioned as a class"
    ],
    "temperature": null,
    "seed": null,
    "weights_hash": null,
    "source_snapshot_id": null,
    "consensus_enabled": false,
    "private_flag": false
  },
  "artifacts": [
    {
      "artifact_id": "a1",
      "modality": "text",
      "role": "object_under_analysis",
      "origin": "operator paste, No Party Politics Issue Map v1",
      "url": "",
      "retrieved_at": "2026-09-16T00:00:00Z",
      "content_hash": null,
      "snapshot_id": null,
      "tier": "T4",
      "tier_rationale": "Self-published analytical document. Rule 2: does not inherit the tier of sources it cites.",
      "extraction": { "method": "direct paste", "extraction_confidence": 1.0, "truncated": false },
      "redirect_chain": [],
      "fetch_status": "not_fetched",
      "provenance": {
        "earliest_known_source": "operator",
        "republication_depth": 0,
        "embedded_credentials": { "exif_present": false, "c2pa_present": false },
        "manipulation_signals": [],
        "clip_bounds": { "clip_start": null, "clip_end": null, "matched_to_official_record": false },
        "link_fidelity": { "checked": true, "post_says": "platform stances as summarized", "source_says": "adopted platform text confirms all 25 tested claims", "mismatch": false }
      }
    },
    {
      "artifact_id": "e1",
      "modality": "document",
      "role": "candidate_evidence",
      "origin": "American Presidency Project",
      "url": "https://www.presidency.ucsb.edu/documents/2024-republican-party-platform",
      "retrieved_at": "2026-09-16T03:27:00Z",
      "content_hash": "sha256:2140dad5fe7d29d0",
      "snapshot_id": "local:npp/platforms/gop.txt",
      "tier": "T1",
      "tier_rationale": "Issuing-body primary document; T1-equivalent for ATTRIBUTIVE claims about its own contents only. T4 for claims about the world. Schema gap, declared.",
      "extraction": { "method": "BeautifulSoup field-docs-content", "extraction_confidence": 0.97, "truncated": false },
      "redirect_chain": [],
      "fetch_status": "held_from_earlier_in_session",
      "provenance": {
        "earliest_known_source": "Republican National Convention, adopted 2024-07-15",
        "republication_depth": 1,
        "embedded_credentials": { "exif_present": false, "c2pa_present": false },
        "manipulation_signals": [
          { "signal": "soft line break splits 'Artificial Intelligence' in archived HTML", "region": "Chapter Three AI subsection heading", "methodology_cited": "regex match failure on contiguous phrase" }
        ]
      }
    }
  ]
}
```

```

    "clip_bounds": { "clip_start": null, "clip_end": null, "matched_to_official_record": false },
    "link_fidelity": { "checked": true, "post_says": "", "source_says": "", "mismatch": false }
  },
  {
    "artifact_id": "e2",
    "modality": "document",
    "role": "candidate_evidence",
    "origin": "American Presidency Project",
    "url": "https://www.presidency.ucsb.edu/documents/2024-democratic-party-platform",
    "retrieved_at": "2026-09-16T03:27:00Z",
    "content_hash": "sha256:55ccfd525108c086",
    "snapshot_id": "local:npp/platforms/dem.txt",
    "tier": "T1",
    "tier_rationale": "Issuing-body primary document; T1-equivalent for ATTRIBUTIVE claims about its own contents
only.",
    "extraction": { "method": "BeautifulSoup field-docs-content", "extraction_confidence": 0.97, "truncated": false },
    "redirect_chain": [],
    "fetch_status": "held_from_earlier_in_session",
    "provenance": {
      "earliest_known_source": "Democratic National Convention, adopted 2024-08-19",
      "republication_depth": 1,
      "embedded_credentials": { "exif_present": false, "c2pa_present": false },
      "manipulation_signals": [],
      "clip_bounds": { "clip_start": null, "clip_end": null, "matched_to_official_record": false },
      "link_fidelity": { "checked": true, "post_says": "", "source_says": "", "mismatch": false }
    }
  },
  "input_is_untrusted": true,
  "injection_detected": false,
  "denatured_claims": [
    "The 2024 Republican and Democratic platforms state the positions this document attributes to them across fifteen
issue areas.",
    "Named nonpartisan pollsters measured the partisan opinion figures reported, on the dates given.",
    "Every line in the document points to a source listed at the bottom.",
    "The named issues will or will not carry into the 2028 presidential cycle."
  ],
  "atomic_claims": [
    { "id": "c1", "text": "The 2024 GOP platform does not mention the Affordable Care Act by name.", "type":
"ATTRIBUTIVE", "subtype": "statutory-content", "load_bearing": true, "resolvable_against": ["T1"], "supporting_artifacts":
["e1"], "verdict": "SUPPORTED", "confidence": 0.95, "evidence_quality": "pinpoint", "model_disagreement": null,
"citation_pending": false, "citations": [ { "url": "https://www.presidency.ucsb.edu/documents/2024-republican-party-
platform", "retrieved_at": "2026-09-16T03:27:00Z", "content_hash": "sha256:2140dad5fe7d29d0", "tier": "T1",
"quoted_passage": "zero occurrences of 'Affordable Care' or 'Obamacare' in full platform text", "extraction_confidence":
0.97 } ] },
    { "id": "c2", "text": "The 2024 GOP platform uses the word 'abortion' once.", "type": "ATTRIBUTIVE", "subtype":
"statutory-content", "load_bearing": false, "resolvable_against": ["T1"], "supporting_artifacts": ["e1"], "verdict":
"SUPPORTED", "confidence": 0.95, "evidence_quality": "pinpoint", "model_disagreement": null, "citation_pending": false,
"citations": [ { "url": "https://www.presidency.ucsb.edu/documents/2024-republican-party-platform", "retrieved_at":
"2026-09-16T03:27:00Z", "content_hash": "sha256:2140dad5fe7d29d0", "tier": "T1", "quoted_passage": "exactly one
occurrence, case-insensitive", "extraction_confidence": 0.97 } ] },
    { "id": "c3", "text": "The GOP platform pledges no change to the Social Security retirement age.", "type":
"ATTRIBUTIVE", "subtype": "statutory-content", "load_bearing": true, "resolvable_against": ["T1"], "supporting_artifacts":
["e1"], "verdict": "SUPPORTED", "confidence": 0.95, "evidence_quality": "pinpoint", "model_disagreement": null,
"citation_pending": false, "citations": [ { "url": "https://www.presidency.ucsb.edu/documents/2024-republican-party-
platform", "retrieved_at": "2026-09-16T03:27:00Z", "content_hash": "sha256:2140dad5fe7d29d0", "tier": "T1",
"quoted_passage": "FIGHT FOR AND PROTECT SOCIAL SECURITY AND MEDICARE WITH NO CUTS, INCLUDING NO CHANGES TO THE RETIREMENT
AGE", "extraction_confidence": 0.97 } ] },
    { "id": "c4", "text": "The GOP platform calls for reasserting federal control over Washington DC.", "type":
"ATTRIBUTIVE", "subtype": "statutory-content", "load_bearing": true, "resolvable_against": ["T1"], "supporting_artifacts":
["e1"], "verdict": "SUPPORTED", "confidence": 0.95, "evidence_quality": "pinpoint", "model_disagreement": null,
"citation_pending": false, "citations": [ { "url": "https://www.presidency.ucsb.edu/documents/2024-republican-party-
platform", "retrieved_at": "2026-09-16T03:27:00Z", "content_hash": "sha256:2140dad5fe7d29d0", "tier": "T1",
"quoted_passage": "Republicans will reassert greater Federal Control over Washington, DC to restore Law and Order in our
Capital City", "extraction_confidence": 0.97 } ] },
    { "id": "c5", "text": "The GOP platform calls for blocking Chinese vehicle imports.", "type": "ATTRIBUTIVE",
"subtype": "statutory-content", "load_bearing": false, "resolvable_against": ["T1"], "supporting_artifacts": ["e1"],
"verdict": "SUPPORTED", "confidence": 0.95, "evidence_quality": "pinpoint", "model_disagreement": null,
"citation_pending": false, "citations": [ { "url": "https://www.presidency.ucsb.edu/documents/2024-republican-party-
platform", "retrieved_at": "2026-09-16T03:27:00Z", "content_hash": "sha256:2140dad5fe7d29d0", "tier": "T1",
"quoted_passage": "preventing the importation of Chinese vehicles", "extraction_confidence": 0.97 } ] },
    { "id": "c6", "text": "The Democratic platform commits to passing the John R. Lewis Voting Rights Advancement Act and
the Freedom to Vote Act.", "type": "ATTRIBUTIVE", "subtype": "statutory-content", "load_bearing": true,
"resolvable_against": ["T1"], "supporting_artifacts": ["e2"], "verdict": "SUPPORTED", "confidence": 0.95,
"evidence_quality": "pinpoint", "model_disagreement": null, "citation_pending": false, "citations": [ { "url": "https://
www.presidency.ucsb.edu/documents/2024-democratic-party-platform", "retrieved_at": "2026-09-16T03:27:00Z", "content_hash":

```

```

"sha256:55ccfd525108c086", "tier": "T1", "quoted_passage": "Democrats will pass and President Biden will sign the John R. Lewis Voting Rights Advancement Act and the Freedom to Vote Act", "extraction_confidence": 0.97 } ] },
{ "id": "c7", "text": "The Democratic platform contains all five gun planks the Issue Map attributes to it.", "type": "ATTRIBUTIVE", "subtype": "statutory-content", "load_bearing": true, "resolvable_against": ["T1"], "supporting_artifacts": ["e2"], "verdict": "SUPPORTED", "confidence": 0.95, "evidence_quality": "pinpoint", "model_disagreement": null, "citation_pending": false, "citations": [ { "url": "https://www.presidency.ucsb.edu/documents/2024-democratic-party-platform", "retrieved_at": "2026-09-16T03:27:00Z", "content_hash": "sha256:55ccfd525108c086", "tier": "T1", "quoted_passage": "Democrats will establish universal background checks... We will once again ban assault weapons and high-capacity magazines. We will require safe storage for guns. Democrats will end the gun industry's immunity from liability... We will pass a national red flag law", "extraction_confidence": 0.97 } ] },
{ "id": "c8", "text": "The Democratic platform's own estimate is that tariffs raise middle-class family taxes by $2,500 a year.", "type": "ATTRIBUTIVE", "subtype": "statistical", "load_bearing": false, "resolvable_against": ["T1"], "supporting_artifacts": ["e2"], "verdict": "SUPPORTED", "confidence": 0.95, "evidence_quality": "pinpoint", "model_disagreement": null, "citation_pending": false, "citations": [ { "url": "https://www.presidency.ucsb.edu/documents/2024-democratic-party-platform", "retrieved_at": "2026-09-16T03:27:00Z", "content_hash": "sha256:55ccfd525108c086", "tier": "T1", "quoted_passage": "tariffs will slow growth, cost jobs, and trigger a recession his first year back in office, raising taxes on middle-class families by $2,500 a year", "extraction_confidence": 0.97 } ] },
{ "id": "c9", "text": "The Democratic platform commits to passing the Equality Act.", "type": "ATTRIBUTIVE", "subtype": "statutory-content", "load_bearing": true, "resolvable_against": ["T1"], "supporting_artifacts": ["e2"], "verdict": "SUPPORTED", "confidence": 0.95, "evidence_quality": "pinpoint", "model_disagreement": null, "citation_pending": false, "citations": [ { "url": "https://www.presidency.ucsb.edu/documents/2024-democratic-party-platform", "retrieved_at": "2026-09-16T03:27:00Z", "content_hash": "sha256:55ccfd525108c086", "tier": "T1", "quoted_passage": "Democrats will pass the Equality Act to codify protections for LGBTQI+ Americans and their families", "extraction_confidence": 0.97 } ] },
{ "id": "c10", "text": "Republicans 61% and Democrats 74% are not confident in the government's ability to regulate AI, per NBC/AIPI.", "type": "EMPIRICAL", "subtype": "statistical", "load_bearing": true, "resolvable_against": ["T2"], "supporting_artifacts": ["a1"], "verdict": "UNSUPPORTED", "confidence": 0.30, "evidence_quality": "attribution disputed", "model_disagreement": null, "citation_pending": true, "citations": [ ] },
{ "id": "c11", "text": "Republicans 98%, independents 84%, Democrats 67% support requiring photo ID to vote.", "type": "EMPIRICAL", "subtype": "statistical", "load_bearing": true, "resolvable_against": ["T2"], "supporting_artifacts": ["a1"], "verdict": "UNSUPPORTED", "confidence": 0.30, "evidence_quality": "sourced to interested-party press release, T4", "model_disagreement": null, "citation_pending": true, "citations": [ ] },
{ "id": "c12", "text": "Border Patrol encounters at the U.S.-Mexico border fell to 237,538 in fiscal 2025, the lowest since 1970.", "type": "EMPIRICAL", "subtype": "statistical", "load_bearing": true, "resolvable_against": ["T1"], "supporting_artifacts": ["a1"], "verdict": "UNSUPPORTED", "confidence": 0.30, "evidence_quality": "metric definition not pinned across comparison span", "model_disagreement": null, "citation_pending": true, "citations": [ ] },
{ "id": "c13", "text": "The average ACA enrollee's premium more than doubled and about 9% of last year's enrollees are now uninsured.", "type": "EMPIRICAL", "subtype": "statistical", "load_bearing": true, "resolvable_against": ["T2"], "supporting_artifacts": ["a1"], "verdict": "UNSUPPORTED", "confidence": 0.30, "evidence_quality": "baseline and denominator unstated", "model_disagreement": null, "citation_pending": true, "citations": [ ] },
{ "id": "c14", "text": "The House passed the SAVE America Act on Feb 11, 2026, requiring photo ID for in-person and mail voting plus DHS citizenship checks of voter rolls.", "type": "LEGAL", "subtype": "vote", "load_bearing": true, "resolvable_against": ["T0", "T1"], "supporting_artifacts": ["a1"], "verdict": "UNSUPPORTED", "confidence": 0.30, "evidence_quality": "no bill number, no Congress number, statutory summary unverified", "model_disagreement": null, "citation_pending": true, "citations": [ ] },
{ "id": "c15", "text": "On Feb 20, 2026 the Supreme Court threw much of the White House's tariff policy into question.", "type": "LEGAL", "subtype": "factual-event", "load_bearing": false, "resolvable_against": ["T0"], "supporting_artifacts": ["a1"], "verdict": "UNSUPPORTED", "confidence": 0.30, "evidence_quality": "not fetched this run", "model_disagreement": null, "citation_pending": true, "citations": [ ] },
{ "id": "c16", "text": "Every line in the document points to a source listed at the bottom.", "type": "EMPIRICAL", "subtype": "provenance", "load_bearing": false, "resolvable_against": ["T4"], "supporting_artifacts": ["a1"], "verdict": "MISLEADING_CONTEXT", "confidence": 0.90, "evidence_quality": "internal inspection", "model_disagreement": null, "citation_pending": false, "citations": [ ] },
{ "id": "c17", "text": "In 2025, 67% of Republicans said the country should prioritize fossil fuels; in 2020, 65% of Republicans prioritized renewables.", "type": "EMPIRICAL", "subtype": "statistical", "load_bearing": true, "resolvable_against": ["T2"], "supporting_artifacts": ["a1"], "verdict": "UNSUPPORTED", "confidence": 0.30, "evidence_quality": "large reversal claim, 2020 figure relayed through 2025 report", "model_disagreement": null, "citation_pending": true, "citations": [ ] },
{ "id": "c18", "text": "The remaining approximately 120 partisan survey figures reported in the document.", "type": "EMPIRICAL", "subtype": "statistical", "load_bearing": true, "resolvable_against": ["T2"], "supporting_artifacts": ["a1"], "verdict": "UNSUPPORTED", "confidence": 0.30, "evidence_quality": "class disposition; fetch disabled; all are retrieval targets", "model_disagreement": null, "citation_pending": true, "citations": [ ] },
{ "id": "c19", "text": "The fifteen '2028 note' assessments of whether each issue carries into the next presidential cycle.", "type": "PREDICTIVE", "subtype": "causal", "load_bearing": false, "resolvable_against": [ ], "supporting_artifacts": ["a1"], "verdict": "UNVERIFIABLE", "confidence": 0.0, "evidence_quality": "predictive; correctly labeled as judgment by the author", "model_disagreement": null, "citation_pending": false, "citations": [ ] } ],
"normative_decomposition": [ ],
"referent_candidates": [
{ "candidate": "The 61/74 AI-confidence pair originates with Pew, Americans and AI 2026, fieldwork 2026-02-17 to 02-23, pp. 17-18, and is relayed by NBC.", "is_hypothesis": true, "confidence": 0.7, "distortion_applied": "pollster label crossed with adjacent AIPI items", "reasoning": "The identical integer pair on the identical item appears in a Pew report already held by this project. Two pollsters returning the same two integers is improbable.", "retrieval_target": "https://www.pewresearch.org/wp-content/uploads/sites/20/2026/06/PI_2026.06.17_Americans-and-AI_REPORT.pdf pp.17-18" },
{ "candidate": "The 237,538 figure is a Border Patrol apprehensions count on a series continuous to 1970, making the comparison valid under a different metric name.", "is_hypothesis": true, "confidence": 0.45, "distortion_applied": "metric renamed 'encounters' in a series where the historical comparison is apprehensions", "reasoning": "'Encounters' entered

```

```

routine publication after 2020 and combines Title 8 apprehensions with expulsions; the underlying apprehensions series is
continuous.", "retrieval_target": "CBP nationwide encounters and apprehensions time series" },
{ "candidate": "The abcnews.com host is a legitimate redirect to abcnews.go.com rather than a lookalike domain.",
"is_hypothesis": true, "confidence": 0.6, "distortion_applied": "none; possible transcription artifact", "reasoning": "ABC
News' canonical host is abcnews.go.com. Cannot be resolved with fetching off.", "retrieval_target": "DNS and redirect
check on abcnews.com" }
},
"conditions_required_for_truth": [
    "Each cited survey page reports the figure attributed to it, for the party group and fieldwork date given.",
    "The 61/74 AI pair is attributable to the pollster named.",
    "The border comparison uses one consistent metric from 1970 to fiscal 2025.",
    "The SAVE America Act text contains no exception that the one-line summary drops."
],
"structural_impossibilities": [],
"plain_language": {
    "gate_triggered": false,
    "pass_a_elements": [],
    "rendering": "",
    "fk_grade": null,
    "modal_verbs_preserved": true,
    "glossary_terms_carried": [],
    "fidelity_score": null,
    "negation_scope_linter": { "source_counts": {}, "render_counts": {}, "mismatch": false },
    "regeneration_attempts": 0,
    "fell_back_to_verbatim": false,
    "pass_c_contaminated": false
},
"domain_failure_patterns": [
    { "pattern": "evidentiary asymmetry by party", "artifact_id": "a1", "span": "Issues 4, 7, 8, 10, 14, 15 Official
stance citations", "mechanism": "Six of fifteen Democratic stances cite third-party summaries while all fifteen Republican
stances cite platform text; the adopted platform contains all six planks.", "severity": "material" },
    { "pattern": "introduced framed as enacted", "artifact_id": "a1", "span": "[DNC_DRAFT] used for Issues 7 and 10",
"mechanism": "A platform draft release is used as the source of record for adopted party positions.", "severity":
"material" },
    { "pattern": "source class mislabeled", "artifact_id": "a1", "span": "[ABC_S0TU24] under 'Democratic Party (2024
platform)'", "mechanism": "A State of the Union liveblog is cited as the party platform.", "severity": "material" },
    { "pattern": "citation-stacking to a T4 origin", "artifact_id": "a1", "span": "Issue 10, 98% / 67% / 84% voter ID
figures", "mechanism": "Survey percentages sourced to an interested party's press release rather than the pollster.",
"severity": "material" },
    { "pattern": "probable source misattribution", "artifact_id": "a1", "span": "Issue 13, '61% not confident' and '74%
not confident' [NBC_AIP126]", "mechanism": "Identical integer pair appears in a Pew report on the same item; advocacy
caveat may be misapplied.", "severity": "material" },
    { "pattern": "precision beyond the data", "artifact_id": "a1", "span": "midterm table Edge column, 'AI ... R+1',
'Economy ... D+1'", "mechanism": "One-point differences ranked as directional with no margin of error published anywhere
in the document.", "severity": "material" },
    { "pattern": "wrong denominator", "artifact_id": "a1", "span": "'237,538 in fiscal 2025, the lowest since 1970'",
"mechanism": "'Encounters' postdates 2020 and combines apprehensions with expulsions; the 1970 comparator is
apprehensions.", "severity": "material" },
    { "pattern": "missing baseline", "artifact_id": "a1", "span": "'the average enrollee's premium more than doubled'",
"mechanism": "Net-of-subsidy and gross premium diverge sharply when a subsidy lapses; the sentence does not say which.",
"severity": "material" },
    { "pattern": "unstated selection rule", "artifact_id": "a1", "span": "'These are the three most relevant measured
positions per group'", "mechanism": "No rule given for 'most relevant'; item selection is unprotected against framing.",
"severity": "material" },
    { "pattern": "header overclaim", "artifact_id": "a1", "span": "'Every line points to a source at the bottom'",
"mechanism": "2028 notes carry no source; Issue 6 cites only [GOP24] for a claim about both platforms.", "severity":
"minor" },
    { "pattern": "compression distortion", "artifact_id": "a1", "span": "'tax incentives for first-time buyers'",
"mechanism": "Platform lists tax incentives and first-time-buyer support as two items; the map merges them.", "severity":
"minor" },
    { "pattern": "stale figure without in-body ageing", "artifact_id": "a1", "span": "Issues 8 and 11 body figures",
"mechanism": "2021, 2023 and 2024 figures appear undated in the body where a reader will quote them.", "severity":
"minor" }
],
"argument_reconstruction": {
    "stated_premises": [
        "Party platforms state official positions.",
        "Nonpartisan surveys measure what each party's voters believe.",
        "The two can be laid side by side to reveal divergence."
    ],
    "unstated_premises": [
        "A platform plank and a survey response distribution are commensurable units.",
        "Three selected survey items fairly represent a party's electorate on an issue.",
        "Differences of one to three points in agreement figures are directional."
    ],
    "conclusion": "Readers can see where each party's stated positions diverge from its own voters' measured positions."
},

```

```

"fallacies": [
  { "name": "base rate / precision neglect", "artifact_id": "a1", "span": "midterm table Edge column", "mechanism":
"Sub-margin differences presented alongside double-digit ones in a single ranked column.", "severity": "material" },
  { "name": "cherry-picking exposure", "artifact_id": "a1", "span": "'three most relevant measured positions per
group'", "mechanism": "Unstated selection rule leaves item choice unconstrained; no slanted instance was found.",
"severity": "material" },
  { "name": "vague quantifier", "artifact_id": "a1", "span": "Issue 9, 'More than 8 in 10' vs '5% to 10%'", "mechanism":
"Two precisions on the two sides of one comparison; range spans two questions without saying so.", "severity": "minor" }
],
"rhetorical_devices": [
  { "device": "smuggled presupposition", "artifact_id": "a1", "span": "Edge column header" }
],
"influence_patterns": [],
"red_team": {
  "counter_evidence": [
    { "claim_id": "F1", "counter": "The adopted Democratic platform is ~43,000 words against the Republican platform's
~5,400 and is largely retrospective, making forward-plank extraction genuinely harder; the asymmetry may be an artifact of
document length rather than intent.", "tier": "T1", "citation_pending": false }
  ],
  "steelman": "The Issue Map is a v1 scoping document that repeatedly and voluntarily discloses its own gaps, flags its
advocacy-aligned pollsters at point of use, declines to fill eleven data holes with plausible numbers, separates a
source's interpretation from platform text unprompted, and attributes contested framing to the party that used it. Judged
as a first pass rather than a publication, its sourcing discipline is above the norm for this genre and the findings here
are refinements, not corrections.",
  "overreach_found": [
    "Initial flag that 'John Lewis Voting Rights Advancement Act' was absent from the adopted platform: false negative
caused by the platform's use of 'John R. Lewis'. Retracted before verdict.",
    "Initial flag that 'reassert federal control over Washington DC' was an interpretation added to a core promise: the
crime chapter states it near-verbatim. Retracted before verdict.",
    "Initial suspicion that this project's prior deliverable misquoted the GOP AI section heading: caused by a soft line
break in the archived HTML ('Artificial Intelligence'). Retracted; prior quote is correct."
  ],
  "tier_confidence_mismatches": [
    "The document flags AIPI and Data for Progress at point of use but flags the White House press release only in a
trailing caveat; the stronger flag belongs on the weaker source."
  ],
  "tier_inflation_found": [
    "[WH_VOTERID] press release treated as a survey source.",
    "[DNC_DRAFT] draft release treated as adopted platform text.",
    "[ABC_50TU24] speech liveblog treated as platform text."
  ],
  "downgrades": [],
  "red_team_contaminated": true
},
"symmetry_check": "Reversed, F1 would fire identically at the same severity with the same fix. The two material sourcing
findings already land on opposite sides: F1 concerns Democratic stances sourced to third parties, F4 concerns a
Republican-administration press release used as a poll source. Neither was softened. F5 concerns a figure on which both
parties look similar, so no partisan gradient is available there.",
"post_level_rollup": { "verdict": "MIXED", "confidence": 0.55, "advisory": true },
"unverified_items": [
  "All partisan survey figures across Issues 1-15 (approximately 120 values).",
  "Midterm focus table, all twelve rows.",
  "Border encounter figure 237,538 and the 'lowest since 1970' comparison.",
  "ACA premium doubling and the 9% uninsured figure.",
  "SAVE America Act passage date, vote, and statutory summary.",
  "February 20, 2026 Supreme Court tariff ruling.",
  "Energy priority reversal, 67% 2025 against 65% 2020.",
  "Voter ID figures 98 / 84 / 67.",
  "AI confidence pair 61 / 74 and its pollster attribution.",
  "abcnews.com host resolution."
],
"unretrievable_artifacts": [
  { "artifact_id": "all 46 reference entries in a1", "reason": "fetch disabled; no --fetch flag passed" }
],
"retrieval_plan": {
  "primary_sources": [
    "Pew, Americans and AI 2026, pp. 17-18",
    "Congress.gov, SAVE America Act with Congress number",
    "clerk.house.gov roll-call record, 2026-02-11",
    "CBP nationwide encounters and apprehensions time series",
    "Gallup, October 2024 voter identification release",
    "KFF, March 2026 tracking poll methodology",
    "Pew methodology pages for every cited survey",
    "Supreme Court opinion, 2026-02-20 tariff decision"
  ],
  "queries": [
    "Pew \"Americans and AI 2026\" \"not confident\" government regulate AI Republicans Democrats",

```

```

    "\"SAVE America Act\" roll call February 11 2026 clerk.house.gov",
    "CBP nationwide Border Patrol apprehensions fiscal year 2025 237,538",
    "Gallup October 2024 voter identification photo ID party Republicans Democrats independents",
    "KFF March 2026 ACA enhanced premium tax credits average premium increase methodology",
    "NBC News June 2026 AI Policy Institute poll regulation confidence"
  ],
  "decisive_artifact": "Pew, Americans and AI 2026 PDF, pp. 17-18. One fetch resolves F5, the only finding that could
mean a published figure is attached to the wrong pollster.",
  "branch_outcomes": [
    { "if_found": "Pew pp. 17-18 carries the 61/74 pair", "then_verdict": "F5 confirmed; relabel to Pew and remove the
advocacy discount on that line" },
    { "if_found": "NBC independently reports 61/74 from AIPI", "then_verdict": "F5 dissolves; note the coincidence and
retain both citations" },
    { "if_found": "CBP series is consistent apprehensions to 1970", "then_verdict": "F7 downgrades to minor; pin the
metric name in the sentence" },
    { "if_found": "CBP 'encounters' series begins post-2020", "then_verdict": "F7 stands; rewrite the comparison or drop
'since 1970'" },
    { "if_found": "Gallup October 2024 shows 98/84/67", "then_verdict": "F4 becomes a citation swap only; numbers
retained" },
    { "if_found": "Gallup shows different figures", "then_verdict": "F4 escalates to disqualifying; remove the three
figures" },
    { "if_found": "SAVE Act text matches the one-line summary", "then_verdict": "Stage 6 gate passes; add bill and
Congress number" },
    { "if_found": "SAVE Act text carries unsummarized exceptions", "then_verdict": "Stage 6 gate fails; render
verbatim" }
  ]
},
"refusals": [],
"handoff_notes": "No artifact markers were supplied; per section 2.2 the entire pasted blob was treated as one artifact
al, modality text, role object_under_analysis, with the preceding line as operator instruction. No --fetch flag was
passed, so the run is pre-retrieval and no claim requiring external retrieval reached SUPPORTED or CONTRADICTED. Two
exceptions are declared and used: the 2024 Republican and Democratic platform texts were retrieved earlier in this same
session from the identical presidency.ucsb.edu URLs the document cites, and are held locally; 25 platform claims were
adjudicated against them. Both UNSUPPORTED and UNVERIFIABLE appear in this run, so the section 5 sanity check passes.
Atomization is a prioritized subset: every claim testable against held evidence, every claim carrying a failure pattern,
and the load-bearing document-level claims; the remaining ~120 survey figures are dispositioned as a class under c18. The
tier schema does not cover party platforms, which are neither law nor official government data nor research; they were
tiered T1-equivalent for attributive claims about their own contents only, on the same principle that makes official
transcripts T1, and T4 for claims about the world. This gap should be closed in v5 with an explicit issuing-body primary-
document tier. Three candidate findings were retracted during the red-team pass after checking the source text; all three
are recorded in red_team.overreach_found rather than dropped, because the retraction rate is itself a result."
}

```